

1. Download all the data in this folder <https://drive.google.com/open?id=1Z4TyI7FcFVEx8qpath/to/the/image.tif,category>

where the categories are numbered 0 to 15, in the following order:

- 0 letter
- 1 form
- 2 email
- 3 handwritten
- 4 advertisement
- 5 scientific report
- 6 scientific publication
- 7 specification
- 8 file folder
- 9 news article
- 10 budget
- 11 invoice
- 12 presentation
- 13 questionnaire
- 14 resume
- 15 memo

2. On this image data, you have to train 3 types of models as given below. You have to s

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

have given th

or you can use this method also

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from->

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, a

5. Use tensorboard for every model and analyse your gradients. (you need to upload the s

Note: fit_generator() method will have problems with the tensorboard histograms, try to

6. You can check about Transfer Learning in this link - <https://blog.keras.io/building-p>

06750000/00484516897554883881/03543900857199698311/1Z4TyI7FcFVEx8qd14j09qxvxaqLSqoEu?e=dow



```
--2020-07-08 11:15:56-- https://doc-0c-0g-docs.googleusercontent.com/docs/securesc/4
Resolving doc-0c-0g-docs.googleusercontent.com (doc-0c-0g-docs.googleusercontent.com)
Connecting to doc-0c-0g-docs.googleusercontent.com (doc-0c-0g-docs.googleusercontent
HTTP request sent, awaiting response... 200 OK
Length: unspecified [application/rar]
Saving to: 'rvl-cdip.rar'
```

```
rvl-cdip.rar          [      <=>          ]  4.34G  80.3MB/s   in 85s
```

```
2020-07-08 11:17:22 (52.1 MB/s) => 'rvl-cdip.rar' saved [4669541700]
```

```
!unrar x '/content/rvl-cdip.rar'
```

```
import pandas as pd
data = pd.read_csv("labels_final.csv")
```

```
data.head()
```

```

┌─┐
└─┘
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
# separating data into train and test.
from sklearn.model_selection import train_test_split
train, test = train_test_split(data, test_size=0.2, random_state=42)
```

```
import tensorflow as tf
import os
import numpy as np
import pandas as pd
```

```
print(train.shape)
print(test.shape)
```

```

┌─┐ (38400, 2)
└─┘ (9600, 2)
```

▼ Model-1

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Max Pooling layer)
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Max Pooling Layer --> FC layers --> Output layer**
4. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

```

from keras import applications
from keras.preprocessing.image import ImageDataGenerator
from keras import optimizers
from keras.models import Sequential
from keras.layers import Dropout, Flatten, Dense
from tensorflow.keras.layers import Dense, Input, Conv2D, MaxPool2D, Activation, Dropout, Flatten
from tensorflow.keras.models import Model
import random as rn
import os
import datetime
datagen=ImageDataGenerator(rescale=1./255)

```

↳ Using TensorFlow backend.

```

train_generator=datagen.flow_from_dataframe(dataframe=train, directory="data_final", x_col="image", y_col="class",
                                             class_mode="classification")
test_generator=datagen.flow_from_dataframe(dataframe=test, directory="/content/data_final", x_col="image", y_col="class",
                                           class_mode="classification")

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

train_size=int(train.shape[0]/32)
test_size = int(test.shape[0]/32)
print(train_size)
print(test_size)

```

↳ 1200
300

```
os.environ['PYTHONHASHSEED'] = '0'
```

```

##https://keras.io/getting-started/faq/#how-can-i-obtain-reproducible-results-using-keras-
## Have to clear the session. If you are not clearing, Graph will create again and again a
## Variables will also set to some value from before session
tf.keras.backend.clear_session()

```

```

## Set the random seed values to regenerate the model.
np.random.seed(0)
rn.seed(0)

```

```

#Input layer
input_layer = Input(shape=(224, 224, 3), name='Input_layer')

```

```

input_layer = Input(shape=(224,224,3),name='input_layer')

#VGG model
vgg_model = tf.keras.applications.VGG16(include_top=False, weights='imagenet',)
print('Model loaded.')
vgg_model.trainable=False
vgg= vgg_model(input_layer)

#Conv Layer
Conv1 = Conv2D(filters=128,kernel_size=(3,3),strides=(1,1),padding='valid',data_format='ch
          activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=0)
#MaxPool Layer
Pool1 = MaxPool2D(pool_size=(2,2),strides=(2,2),padding='valid',data_format='channels_last

#Flatten
flatten = Flatten(data_format='channels_last',name='Flatten')(Pool1)

#FC layer
FC1 = Dense(units=64,activation='relu',kernel_initializer=tf.keras.initializers.glorot_nor

#FC layer
FC2 = Dense(units=32,activation='relu',kernel_initializer=tf.keras.initializers.glorot_nor

#output layer
Out = Dense(units=16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_

#Creating a model
model1= Model(inputs=input_layer,outputs=Out)

```

Automatic saving failed. This file was updated remotely or in another tab.

[Show](#)

[pplications/vg](#)

[diff](#)

Model loaded.

```
model1.summary()
```



Model: "model"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```
#compiling
model1.compile(optimizer=tf.keras.optimizers.Adam(lr=0.01),loss='categorical_crossentropy',
               metrics=['accuracy'],mode='train')

%load_ext tensorboard
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)

#tensorboard --logdir $logdir
```



Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

TensorBoard

SCALARS

GRAPHS

INACTIVE

```
model1.fit_generator(generator=train_generator,steps_per_epoch=train_size,validation_data=
```

```

[ ] WARNING:tensorflow:From <ipython-input-17-e882bde338b6>:1: Model.fit_generator (from
Instructions for updating:
Please use Model.fit, which supports generators.
Epoch 1/15
1200/1200 [=====] - 296s 247ms/step - loss: 2.0599 - accurac
Epoch 2/15
1200/1200 [=====] - 289s 241ms/step - loss: 1.6197 - accurac
Epoch 3/15
1200/1200 [=====] - 278s 232ms/step - loss: 1.4799 - accurac
Epoch 4/15
1200/1200 [=====] - 274s 229ms/step - loss: 1.4088 - accurac
Epoch 5/15
1200/1200 [=====] - 272s 227ms/step - loss: 1.3690 - accurac
Epoch 6/15
1200/1200 [=====] - 270s 225ms/step - loss: 1.3351 - accurac
Epoch 7/15
1200/1200 [=====] - 271s 226ms/step - loss: 1.3008 - accurac
Epoch 8/15
1200/1200 [=====] - 268s 224ms/step - loss: 1.2787 - accurac
Epoch 9/15
1200/1200 [=====] - 270s 225ms/step - loss: 1.2593 - accurac
Epoch 10/15
1200/1200 [=====] - 270s 225ms/step - loss: 1.2400 - accurac
Epoch 11/15
1200/1200 [=====] - 274s 228ms/step - loss: 1.2106 - accurac
Epoch 12/15
1200/1200 [=====] - 272s 228ms/step - loss: 1.1952 - accurac
Epoch 13/15
1200/1200 [=====] - 273s 228ms/step - loss: 1.1876 - accurac
Epoch 14/15
1200/1200 [=====] - 259s 216ms/step - loss: 1.1830 - accurac
Epoch 15/15
1200/1200 [=====] - 265s 221ms/step - loss: 1.1563 - accurac
<tensorflow.python.keras.callbacks.History at 0x7fcde3f92e80>

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

▼ Model-2

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initilize all the we
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers ident
3. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the

```

tf.keras.backend.clear_session()
## Set the random seed values to regenerate the model.
np.random.seed(0)
rn.seed(0)

```

```

#Input layer
input_layer = Input(shape=(224,224,3,))

#VGG model
vgg_model = tf.keras.applications.VGG16(include_top=False, weights='imagenet',input_shape=
vgg_model.trainable=False

vgg= vgg_model(input_layer)
#Conv Layer
FCConv1 = Conv2D(filters=4096,kernel_size=(7,7),strides=(1,1),padding='valid',data_format=
activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=0)

FCConv2 = Conv2D(filters=4096,kernel_size=(1,1),strides=(1,1),padding='valid',data_format=
activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=30)
#Flatten
flatten = Flatten(data_format='channels_last',name='Flatten')(FCConv2)
#output layer
Out = Dense(units=16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_

#Creating a model
model2_new= Model(inputs=input_layer,outputs=Out)

model2_new.summary()

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

=====
input_1 (InputLayer)      [(None, 224, 224, 3)]    0
-----
vgg16 (Model)             (None, 7, 7, 512)       14714688
-----
conv2d (Conv2D)           (None, 1, 1, 4096)      102764544
-----
conv2d_1 (Conv2D)         (None, 1, 1, 4096)      16781312
-----
Flatten (Flatten)         (None, 4096)            0
-----
Output (Dense)            (None, 16)              65552
=====
Total params: 134,326,096
Trainable params: 119,611,408
Non-trainable params: 14,714,688
=====

```

```

#compiling
model2_new.compile(optimizer=tf.keras.optimizers.Adam(lr=0.001),loss='categorical_crossentropy')

%load_ext tensorboard
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)

```

↗ The tensorboard extension is already loaded. To reload it, use:
`%reload_ext tensorboard`

`%tensorboard --logdir $logdir`



TensorBoard

SCALARS

GRAPHS

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting method: default ▼

Smoothing



0.6

Horizontal Axis

STEP


RELATIVE

WALL

epoch_accuracy



epoch_accuracy



Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



Write a regex to filter runs


☐ ○ train

☐ ○ validation

TOGGLE ALL RUNS

logs/20200708-142917

epoch_loss



`model2_new.fit_generator(generator=train_generator, steps_per_epoch=train_size, validation_d`




```

Epoch 1/15
1200/1200 [=====] - 425s 354ms/step - loss: 1.4449 - accurac
Epoch 2/15
1200/1200 [=====] - 425s 354ms/step - loss: 0.9344 - accurac
Epoch 3/15
1200/1200 [=====] - 425s 354ms/step - loss: 0.7969 - accurac
Epoch 4/15
1200/1200 [=====] - 425s 354ms/step - loss: 0.6980 - accurac
Epoch 5/15
1200/1200 [=====] - 425s 354ms/step - loss: 0.6260 - accurac
Epoch 6/15
1200/1200 [=====] - 425s 354ms/step - loss: 0.5546 - accurac
Epoch 7/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.5096 - accurac
Epoch 8/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.4589 - accurac
Epoch 9/15
1200/1200 [=====] - 427s 356ms/step - loss: 0.4179 - accurac
Epoch 10/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.3950 - accurac
Epoch 11/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.3611 - accurac
Epoch 12/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.3401 - accurac
Epoch 13/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.3119 - accurac
Epoch 14/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.2984 - accurac
Epoch 15/15
1200/1200 [=====] - 426s 355ms/step - loss: 0.2984 - accurac

```

▼ Model-3

Automatic saving failed. This file was updated remotely or in another tab. [Show](#)

[diff](#)

1. Use same network as Model-2 INPUT --> VGG-16 without top layers(FC) --> 2 Conv Layer

```

tf.keras.backend.clear_session()
## Set the random seed values to regenerate the model.
np.random.seed(0)
rn.seed(0)
#Input layer
input_layer = Input(shape=(224,224,3,))

#VGG model
vgg_model = tf.keras.applications.VGG16(include_top=False, weights='imagenet',input_shape=
for i in range(0,19):
    if i<13:
        vgg_model.layers[i].trainable=False
    else:
        vgg_model.layers[i].trainable=True

```

```

vgg= vgg_model(input_layer)
#Conv Layer
FCConv1 = Conv2D(filters=128,kernel_size=(7,7),strides=(1,1),padding='valid',data_format='
            activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=0)

FCConv2 = Conv2D(filters=64,kernel_size=(1,1),strides=(1,1),padding='valid',data_format='c
            activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=30
#Flatten
flatten = Flatten(data_format='channels_last',name='Flatten')(FCConv2)
#output layer
Out = Dense(units=16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_

#Creating a model
model3_new= Model(inputs=input_layer,outputs=Out)

```

```
model3_new.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
vgg16 (Model)	(None, 7, 7, 512)	14714688
conv2d (Conv2D)	(None, 1, 1, 128)	3211392
conv2d_1 (Conv2D)	(None, 1, 1, 64)	8256

Automatic saving failed. This file was updated remotely or in another tab. [Show](#)

[diff](#)

```

output (Dense) (None, 16) 256
=====
Total params: 17,935,376
Trainable params: 12,659,920
Non-trainable params: 5,275,456

```

```

#compiling
model3_new.compile(optimizer=tf.keras.optimizers.Adam(lr=0.001),loss='categorical_crossentropy

```

```

%load_ext tensorboard
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)
%tensorboard --logdir $logdir

```

↗

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

TensorBoard

SCALARS

GRAPHS

INACTIVE

☐ Show data download links

☐ Ignore outliers in chart scaling

Tooltip sorting method:

default

Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

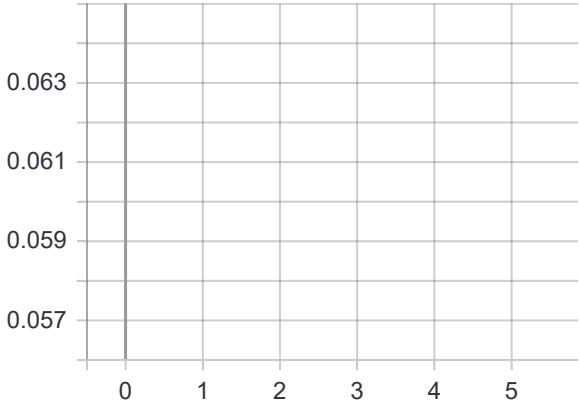
☐ ☐ validation

TOGGLE ALL RUNS

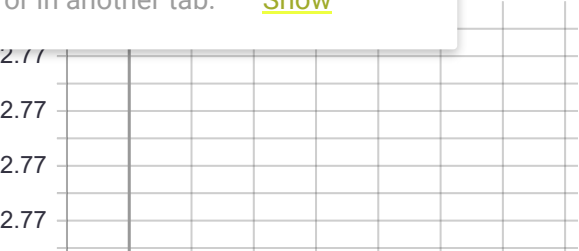
logs/20200708-124823

epoch_accuracy

epoch_accuracy



epoch_loss



```
model3_new.fit_generator(generator=train_generator,steps_per_epoch=train_size,validation_d
```



```
Epoch 1/15
1200/1200 [=====] - 294s 245ms/step - loss: 2.7774 - accurac
Epoch 2/15
1200/1200 [=====] - 292s 243ms/step - loss: 2.7728 - accurac
Epoch 3/15
1200/1200 [=====] - 292s 244ms/step - loss: 2.7728 - accurac
Epoch 4/15
1200/1200 [=====] - 294s 245ms/step - loss: 2.7728 - accurac
Epoch 5/15
1200/1200 [=====] - 290s 242ms/step - loss: 2.7728 - accurac
Epoch 6/15
1200/1200 [=====] - 293s 245ms/step - loss: 2.7728 - accurac
Epoch 7/15
1200/1200 [=====] - 290s 242ms/step - loss: 2.7728 - accurac
Epoch 8/15
1200/1200 [=====] - 293s 244ms/step - loss: 2.7728 - accurac
Epoch 9/15
1200/1200 [=====] - 290s 242ms/step - loss: 2.7728 - accurac
Epoch 10/15
1200/1200 [=====] - 293s 244ms/step - loss: 2.7728 - accurac
Epoch 11/15
1200/1200 [=====] - 290s 242ms/step - loss: 2.7728 - accurac
Epoch 12/15
1200/1200 [=====] - 293s 244ms/step - loss: 2.7728 - accurac
Epoch 13/15
1200/1200 [=====] - 290s 242ms/step - loss: 2.7728 - accurac
Epoch 14/15
1200/1200 [=====] - 292s 243ms/step - loss: 2.7728 - accurac
Epoch 15/15
1200/1200 [=====] - 290s 242ms/step - loss: 2.7728 - accurac
<tensorflow.python.keras.callbacks.History at 0x7fcdd838a860>
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)